

```

-- Your Name: Chuqi Wang (79167724)
-- 1. Table Creation and Analysis --
-- SQL DDL statements
\i /Users/chuqi wang/Desktop/UCI/CS224P/assignments/HW1/ZotMusicDDL.sql
DROP SCHEMA IF EXISTS ZotMusic CASCADE;
DROP SCHEMA
CREATE SCHEMA ZotMusic;
CREATE SCHEMA
SET search_path TO ZotMusic;
SET

CREATE TABLE Users (
    user_id      text NOT NULL,
    email        text NOT NULL,
    joined_date   date NOT NULL,
    nickname     text NOT NULL,
    street        text,
    city          text,
    state         text,
    zip          text,
    genres        text,
    PRIMARY KEY (user_id)
);
CREATE TABLE

CREATE TABLE Artists (
    user_id      text,
    bio          text,
    stagename    text,
    PRIMARY KEY (user_id),
    FOREIGN KEY (user_id) REFERENCES Users (user_id) ON DELETE CASCADE
);
CREATE TABLE

CREATE TABLE Listeners (
    user_id      text,
    subscription text,
    first_name    text NOT NULL,
    last_name     text NOT NULL,
    PRIMARY KEY (user_id),
    FOREIGN KEY (user_id) REFERENCES Users (user_id) ON DELETE
CASCADE,
    CHECK (subscription IN ('free', 'monthly', 'yearly'))
);
CREATE TABLE

CREATE TABLE Records (
    record_id     text NOT NULL,
    artist_user_id text NOT NULL,
    title         text NOT NULL,

```

```

        genre            text NOT NULL,
        release_date     date NOT NULL,
        PRIMARY KEY (record_id),
        FOREIGN KEY (artist_user_id) REFERENCES Artists (user_id) ON
DELETE CASCADE
);
CREATE TABLE

CREATE TABLE Singles (
    record_id            text NOT NULL,
    video_url            text,
    PRIMARY KEY (record_id),
    FOREIGN KEY (record_id) REFERENCES Records (record_id) ON DELETE
CASCADE
);
CREATE TABLE

CREATE TABLE Albums (
    record_id            text NOT NULL,
    description          text,
    PRIMARY KEY (record_id),
    FOREIGN KEY (record_id) REFERENCES Records (record_id) ON DELETE
CASCADE
);
CREATE TABLE

CREATE TABLE Songs (
    record_id            text NOT NULL,
    track_number         int NOT NULL,
    title                text NOT NULL,
    length               int NOT NULL,
    bpm                  int,
    mood                 text,
    PRIMARY KEY (record_id, track_number),
    FOREIGN KEY (record_id) REFERENCES Records (record_id) ON DELETE
CASCADE
);
CREATE TABLE

CREATE TABLE Sessions (
    session_id           text NOT NULL,
    user_id              text NOT NULL,
    record_id            text NOT NULL,
    track_number         int NOT NULL,
    initiate_at          timestamp NOT NULL,
    leave_at             timestamp NOT NULL,
    music_quality        text NOT NULL,
    device               text NOT NULL,
    remaining_time       int NOT NULL,
    replay_count         int,

```

```

        PRIMARY KEY (session_id),
        FOREIGN KEY (user_id) REFERENCES Listeners(user_id) ON DELETE
CASCADE,
        FOREIGN KEY (record_id, track_number) REFERENCES Songs(record_id,
track_number) ON DELETE CASCADE
);

```

CREATE TABLE

```

CREATE TABLE Reviews (
    review_id      text NOT NULL,
    user_id        text NOT NULL,
    record_id      text NOT NULL,
    rating         int NOT NULL,
    body           text,
    posted_at      timestamp NOT NULL,
    PRIMARY KEY (review_id),
    FOREIGN KEY (user_id) REFERENCES Listeners (user_id) ON DELETE
CASCADE,
    FOREIGN KEY (record_id) REFERENCES Records (record_id) ON DELETE
CASCADE
);

```

CREATE TABLE

```

CREATE TABLE ReviewLikes(
    user_id text NOT NULL,
    review_id text NOT NULL,
    PRIMARY KEY (user_id, review_id),
    FOREIGN KEY (user_id) REFERENCES Listeners(user_id) ON DELETE
CASCADE,
    FOREIGN KEY (review_id) REFERENCES Reviews(review_id) ON DELETE
CASCADE
);

```

CREATE TABLE

-- Songs table field design observations

\d songs

Table "zotmusic.songs"				
Column	Type	Collation	Nullable	Default
record_id	text		not null	
track_number	integer		not null	
title	text		not null	
length	integer		not null	
bpm	integer			
mood	text			

Indexes:

"songs\_pkey" PRIMARY KEY, btree (record\_id, track\_number)

Foreign-key constraints:

"songs\_record\_id\_fkey" FOREIGN KEY (record\_id) REFERENCES  
records(record\_id) ON DELETE CASCADE

Referenced by:

```
TABLE "sessions" CONSTRAINT "sessions_record_id_track_number_fkey"  
FOREIGN KEY (record_id, track_number) REFERENCES songs(record_id,  
track_number) ON DELETE CASCADE
```

```
/*  
record_id: text NOT NULL  
track_number: int NOT NULL  
title: text NOT NULL  
length: int NOT NULL  
bpm: int  
mood: text  
*/  
-- 2. Data Loading (COPY commands) --  
\copy Users from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/HW1/  
zot-music-dataset-assignment1/Users.csv' delimiter ',' csv header;  
COPY 200  
\copy Artists from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Artists.csv' delimiter ',' csv  
header;  
COPY 116  
\copy Listeners from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Listeners.csv' delimiter ',' csv  
header;  
COPY 184  
\copy Records from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Records.csv' delimiter ',' csv  
header;  
COPY 1000  
\copy Singles from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Singles.csv' delimiter ',' csv  
header;  
COPY 300  
\copy Albums from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Albums.csv' delimiter ',' csv  
header;  
COPY 700  
\copy Songs from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/HW1/  
zot-music-dataset-assignment1/Songs.csv' delimiter ',' csv header;  
COPY 6252  
\copy Sessions from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Sessions.csv' delimiter ',' csv  
header;  
COPY 50000  
\copy Reviews from '/Users/chuqiwang/Desktop/UCI/CS224P/assignments/  
HW1/zot-music-dataset-assignment1/Reviews.csv' delimiter ',' csv  
header;  
COPY 9499  
\copy ReviewLikes from '/Users/chuqiwang/Desktop/UCI/CS224P/  
assignments/HW1/zot-music-dataset-assignment1/ReviewLikes.csv'  
delimiter ',' csv header;
```

COPY 91325

-- 3. Query Answers --

set search\_path to zotmusic;

SET

-- Problem A --

```
select 'Users' as entity, count(*) from Users union all
select 'Records' as entity, count(*) from Records union all
select 'Reviews' as entity, count(*) from Reviews;
```

entity	count
Users	200
Records	1000
Reviews	9499

(3 rows)

-- Problem B --

```
select u.user_id, u.email, u.nickname, u.zip from users as u
join artists on artists.user_id = u.user_id
join listeners on listeners.user_id = u.user_id
where email like '%@icloud.com';
```

user_id	email
nickname	zip
user_10dfa3b6-52b6-43a9-835a-ad110ad50ff2	roberthammond@icloud.com
roberthammond	54869
user_4a9ffbf6-5430-45a4-b68e-0ae6f6737d8b	william09@icloud.com
william09	55308
user_457ad608-9661-4384-9919-1d89c52fd0de	danielharrison@icloud.com
danielharrison	23703
user_d383327a-dd9c-4a4f-bbaf-262c7a6d90a0	chelsealawson@icloud.com
chelsealawson	10206
user_83a40c0c-573e-44ec-8cac-b5951513f88b	turnerkayla@icloud.com
turnerkayla	02182
user_38eaa9f8-e8fc-4ce4-a8ae-ffb882c1786c	ryanmorgan@icloud.com
ryanmorgan	95166
user_3c5d30bc-0ac2-4df1-8574-892d2f666df6	browncarrie@icloud.com
browncarrie	23550

(7 rows)

-- Problem C --

```
select record_id, title, genre, release_date from records
where artist_user_id = (
    select a.user_id from artists as a
    join users as u on a.user_id = u.user_id
    where email = 'fwilson@outlook.com'
)
```

order by release\_date;

record_id	title
genre	release_date

```

-----
+-----+-----+-----
record_91c6325d-b17f-4f4c-be6d-3517b2173a9f | Statement matter
| Country | 2020-01-12
record_822961a3-946a-49ff-8173-74d4035286b9 | Apply size
| Gospel | 2020-01-29
record_cbf93efd-2deb-48ae-ad73-83aa088c6f13 | Democratic what
| Soul | 2020-03-27
record_2406e933-23e3-4db1-acf9-3c863d48bff6 | General job heavy
| Country | 2020-05-08
record_57061d35-de20-4bf1-9aac-a689f0db7e16 | Would determine
| Soul | 2020-06-07
record_3e4ed054-cf1a-4a04-8e97-e0177c6d3575 | Summer civil political
beat | Folk | 2021-03-31
record_116fbdd6-e706-41f7-9809-12e174e48e8f | Discover rate
| Jazz | 2021-09-09
record_62389a63-e95f-43d1-acea-aalbac0e0050 | Result guess for
| Gospel | 2021-10-17
record_5cbf14c7-7b54-4e32-bfce-cba507c7277f | Bar talk long
| Jazz | 2021-10-23
(9 rows)

```

-- Problem D --

```

with cte as (
    select * from records
    where artist_user_id = (
        select a.user_id from artists as a
        join users as u on a.user_id = u.user_id
        where email = 'fwilson@outlook.com'
    )
)
select artist_user_id, genre,
       count(case when a.record_id is not null then 1 end) as
album_count,
       count(case when s.record_id is not null then 1 end) as
single_count
from cte
left join singles as s on cte.record_id = s.record_id
left join albums as a on cte.record_id = a.record_id
group by artist_user_id, genre;

```

	artist_user_id	genre	album_count	single_count
1	user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3	Country	1	
0	user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3	Folk	1	
0	user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3	Gospel	2	

```

user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3 | Jazz      | 1 |
1
user_6ac27408-a0a6-4c57-a025-7b6854f7a8e3 | Soul      | 1 |
1
(5 rows)

```

-- Problem E --

```

select coalesce(a.stagename, u.nickname) as name,
u.email from users as u
join artists as a on u.user_id = a.user_id
join records as r on u.user_id = r.artist_user_id
group by a.user_id, a.stagename, u.email, u.nickname
having count(distinct r.genre) >= 9;

```

name	email
blakeshannon	william09@icloud.com
elizabeth55	powerschristopher@foxmail.com
ehester	william57@mail.com
richardsbilly	robert92@outlook.com
sandersallison	keith65@university.edu
khall	lthompson@college.edu

(6 rows)

-- Problem F --

```

with cte as (
    select artist_user_id from records as r
    where r.genre in ('R&B', 'Hip-Hop')
    group by artist_user_id
    having count(distinct case when r.genre = 'R&B' then 1 end) >
0 and
    count(distinct case when r.genre = 'Hip-Hop' then 1 end) > 0
)
select u.user_id, u.email from users as u
join cte on u.user_id = cte.artist_user_id
where u.user_id not in (
    select distinct artist_user_id from records as r
    where r.genre in ('Indie', 'Jazz')
)
order by u.user_id;

```

user_id	email
user_377075cd-a1c3-4f19-816c-0fdcf9c973607	cflores@university.edu
user_39c8d999-40e5-4fd0-9cd5-96a93239abee	charleslewis@university.edu
user_3fbded9d-d59a-435c-863f-55f0b086f01e	zhill@hotmail.com
user_5a61e935-0fec-4c35-9bc4-58f9e7ecb067	molly19@protonmail.com
user_80b63994-e7f9-4f57-ab42-88683a70d183	william57@mail.com
user_81b9d067-7877-43cf-9793-6fb5ea6f4921	melissa63@university.edu
user_83a40c0c-573e-44ec-8cac-b5951513f88b	turnerkayla@icloud.com

```

user_94b498aa-2a93-4e8e-8efa-229604beea67 | mcguiresheila@outlook.com
user_b26fbc8f-e3c8-426f-9d4c-77eec730bacd | anthony59@college.edu
user_bf4b7630-b9dd-40f0-b534-ef841ea43194 | linda38@outlook.com
user_f6cb31b3-48fa-4cd1-a483-7a59189af5d1 | joseph37@yahoo.com
(11 rows)

```

-- Problem G --

```

select email, nickname, array_length(string_to_array(genres, ','), 1)
as num_genres
from users
order by num_genres desc
limit 10;

```

email	nickname	num_genres
courtney36@protonmail.com	courtney36	10
charleslewis@university.edu	charleslewis	10
ewilliams@mail.com	ewilliams	10
zmason@gmail.com	zmason	10
gomezbrittany@foxmail.com	gomezbrittany	9
bknapp@icloud.com	bknapp	9
ryanmorgan@icloud.com	ryanmorgan	9
edwardscindy@foxmail.com	edwardscindy	8
joel00@gmail.com	joel00	8
gclayton@protonmail.com	gclayton	8

(10 rows)

-- Problem H --

```

select unnest(string_to_array(genres, ',')) as genre, count(user_id)
as num_users
from users
group by genre
order by num_users desc;

```

genre	num_users
Soul	58
Techno	56
Indie	54
Folk	53
Blues	53
Funk	49
Country	49
Classical	47
R&B	47
Metal	47
Jazz	47
Disco	46
Hip-Hop	45
Latin	45
Pop	44
Gospel	43



Punk		42
Electronic		38
Rock		33
Reggae		27

(20 rows)

-- Problem I --

```
select r.title, r.record_id, s.track_number, s.title,
count(se.user_id) as num_listeners
from sessions as se
join songs as s on se.record_id = s.record_id and se.track_number =
s.track_number
join records as r on se.record_id = r.record_id
where se.remaining_time <= s.length * 0.2 or se.replay_count > 0
group by r.record_id, s.track_number, s.title
order by num_listeners desc
limit 10;
```

title	record_id
track_number   title	num_listeners
-----	
+-----+-----	
+-----+-----	
Small century stop be66-41ddf41245a7	record_154ce2a2-1e2c-48c4- 8   Local stay keep
16	
Popular f1f0731ce72c	record_7e742b33-f00b-4356-83a0- 5   Brother despite   16
Personal do physical c1cf3013fa42	record_3a1932fb-22f1-4834-b0fb- 2   Identify community   16
Benefit him these ba72-4d34-9d81-021c01c2e3fa	record_a120d09b- 1   Police will human
16	
Dog bac88d34f4fb	record_422a9857-f805-4e65-b4f0- 5   Water   15
Understand yourself suggest fa30-471c-8711-128a4e8ff977	record_e5260c91- 5   Direction air
15	
Western fear ae62-00580338f467	record_b6f51c72-3534-4b29- 6   Myself letter across
14	
Stop watch different ab21-3a42beb4987b	record_d9120d24-b69a-41bf- 3   Skill throw
14	
Western parent firm speech c6e900aa2088	record_0104059c-6029-42d6-a1e5- 10   Market test could   14
Such view economy b43b-409fd9fe195c	record_489ad7e4-9a03-44fa- 8   Industry when
14	

(10 rows)

```

-- Problem J --
-- View DDL:
create view RatedRecords as
with cte as (
    select rv.review_id, rv.record_id, rv.rating,
           coalesce(count(rl.user_id), 0) + 1 as weight
    from reviews as rv
    left join reviewlikes as rl on rv.review_id = rl.review_id
    group by rv.review_id, rv.record_id, rv.rating
)
select r.title, r.record_id,
       coalesce(sum(cte.rating * cte.weight) / nullif(sum(cte.weight), 0), 0)
as rating,
       count(cte.review_id) as num_reviews
from records as r
join cte on r.record_id = cte.record_id
group by r.title, r.record_id;
CREATE VIEW
-- View test query:
select * from RatedRecords
where num_reviews >= 5
order by rating desc
limit 10;

```

title	record_id
rating	num_reviews
Second concern star	record_ab8ecafc-cae3-4394-824b-45f860d63c8a
4.3287671232876712	9
Real growth	record_dd3f83b0-f0ec-41a0-b048-83a7dd515191
4.2375000000000000	7
Crime	record_602a6db2-1784-45d2-8645-1c141f25a049
4.2142857142857143	6
News explain might turn	record_cbf50912-23ca-4de0-b94d-41b2076fd9d3
4.2121212121212121	6
Black movement	record_acb14a89-3854-45de-9df7-ef1be87c7e47
4.1724137931034483	5
Water film	record_0e4952cc-3445-4fe6-80c9-867e98ae4233
4.1630434782608696	8
Be sign hair	record_e8fa438d-3e01-40df-bd1c-5d7dadf98015
4.1532258064516129	10
Music religious charge	record_3cd5b91e-f6c5-469a-8f92-98f85a1a7af7
4.1392405063291139	7
Change adult	record_9388a7fd-5828-4438-82b1-46c13e63c6c2
4.1241830065359477	17
Sit pay political	record_66f81619-1bf3-4784-bd94-9c3adeeb5785
4.1086956521739130	6

(10 rows)

-- Problem K --

-- Table alteration DDL:

alter table records

add rating decimal(3, 2);

ALTER TABLE

-- Table update query:

update records as r

set rating = RatedRecords.rating

from RatedRecords

where r.record\_id = RatedRecords.record\_id;

UPDATE 1000

-- Change verification query:

select r.title, r.record\_id, r.rating, RatedRecords.num\_reviews

from records as r

join RatedRecords on r.record\_id = RatedRecords.record\_id

where RatedRecords.num\_reviews >= 5

order by r.rating desc

limit 10;

title	record_id
rating   num_reviews	
Second concern star	record_ab8ecafc-cae3-4394-824b-45f860d63c8a
4.33   9	
Real growth	record_dd3f83b0-f0ec-41a0-b048-83a7dd515191
4.24   7	
Crime	record_602a6db2-1784-45d2-8645-1c141f25a049
4.21   6	
News explain might turn	record_cbf50912-23ca-4de0-b94d-41b2076fd9d3
4.21   6	
Black movement	record_acb14a89-3854-45de-9df7-ef1be87c7e47
4.17   5	
Water film	record_0e4952cc-3445-4fe6-80c9-867e98ae4233
4.16   8	
Be sign hair	record_e8fa438d-3e01-40df-bd1c-5d7dadf98015
4.15   10	
Music religious charge	record_3cd5b91e-f6c5-469a-8f92-98f85a1a7af7
4.14   7	
Change adult	record_9388a7fd-5828-4438-82b1-46c13e63c6c2
4.12   17	
Sit pay political	record_66f81619-1bf3-4784-bd94-9c3adeeb5785
4.11   6	

(10 rows)

-- Problem L --

-- Query against view:

select a.user\_id, u.nickname, avg(rr.rating) as content\_rating

from artists as a

join users as u on a.user\_id = u.user\_id

join records as r on a.user\_id = r.artist\_user\_id

```

join RatedRecords as rr ON r.record_id = rr.record_id
group by a.user_id, u.nickname
having avg(rr.rating) >= 3.3;

```

user_id	nickname	content_rating
user_f921401a-7991-4db2-9491-2cb32b4146db	paynedavid	3.4573236606157646
user_827369a9-7b0e-4937-917e-632d1ed5620f	robertfigueroa	3.3393895455471657

(2 rows)

```

EXPLAIN ANALYZE
select a.user_id, u.nickname, avg(rr.rating) as content_rating
from artists as a
join users as u on a.user_id = u.user_id
join records as r on a.user_id = r.artist_user_id
join RatedRecords as rr ON r.record_id = rr.record_id
group by a.user_id, u.nickname
having avg(rr.rating) >= 3.3;

```

#### QUERY PLAN

```

-----
HashAggregate  (cost=4643.62..4669.70 rows=580 width=96) (actual
time=27.754..27.779 rows=2 loops=1)
  Group Key: a.user_id, u.nickname
  Filter: (avg(((COALESCE((sum((rv.rating *
((COALESCE(count(rl.user_id), '0'::bigint) + 1)))) /
NULLIF(sum(((COALESCE(count(rl.user_id), '0'::bigint) + 1))),
'0'::numeric)), '0'::numeric))) >= 3.3)
  Batches: 1  Memory Usage: 129kB
  Rows Removed by Filter: 114
  -> Hash Join  (cost=4564.63..4630.58 rows=1739 width=96) (actual
time=27.044..27.587 rows=1000 loops=1)
    Hash Cond: (a.user_id = u.user_id)
    -> Hash Join  (cost=4548.33..4609.68 rows=1739 width=96)
(actual time=27.012..27.417 rows=1000 loops=1)
      Hash Cond: (r.artist_user_id = a.user_id)
      -> Hash Join  (cost=4523.71..4580.45 rows=1739
width=64) (actual time=26.994..27.292 rows=1000 loops=1)
        Hash Cond: (r_1.record_id = r.record_id)
        -> HashAggregate  (cost=4447.58..4482.36
rows=1739 width=104) (actual time=26.869..27.044 rows=1000 loops=1)
          Group Key: r_1.record_id
          Batches: 1  Memory Usage: 577kB
          -> Hash Join  (cost=3988.90..4317.02
rows=13056 width=76) (actual time=24.060..25.728 rows=9499 loops=1)

```

```

Hash Cond: (rv.record_id =
r_1.record_id)
-> HashAggregate
(cost=3912.77..4075.97 rows=13056 width=76) (actual
time=23.928..24.547 rows=9499 loops=1)
Group Key: rv.review_id
Batches: 1 Memory Usage:
1937kB
-> Hash Right Join
(cost=565.76..3335.49 rows=115456 width=100) (actual
time=1.383..14.768 rows=91523 loops=1)
Hash Cond: (rl.review_id
= rv.review_id)
-> Seq Scan on
reviewlikes rl (cost=0.00..2466.56 rows=115456 width=64) (actual
time=0.002..2.931 rows=91325 loops=1)
-> Hash
(cost=402.56..402.56 rows=13056 width=68) (actual time=1.372..1.372
rows=9499 loops=1)
Buckets: 16384
Batches: 1 Memory Usage: 1279kB
-> Seq Scan on
reviews rv (cost=0.00..402.56 rows=13056 width=68) (actual
time=0.002..0.602 rows=9499 loops=1)
-> Hash (cost=54.39..54.39
rows=1739 width=64) (actual time=0.129..0.129 rows=1000 loops=1)
Buckets: 2048 Batches: 1
Memory Usage: 107kB
-> Seq Scan on records r_1
(cost=0.00..54.39 rows=1739 width=64) (actual time=0.003..0.050
rows=1000 loops=1)
-> Hash (cost=54.39..54.39 rows=1739 width=64)
(actual time=0.123..0.123 rows=1000 loops=1)
Buckets: 2048 Batches: 1 Memory Usage:
132kB
-> Seq Scan on records r
(cost=0.00..54.39 rows=1739 width=64) (actual time=0.007..0.058
rows=1000 loops=1)
-> Hash (cost=16.50..16.50 rows=650 width=32) (actual
time=0.015..0.015 rows=116 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 17kB
-> Seq Scan on artists a (cost=0.00..16.50
rows=650 width=32) (actual time=0.002..0.007 rows=116 loops=1)
-> Hash (cost=12.80..12.80 rows=280 width=64) (actual
time=0.031..0.031 rows=200 loops=1)
Buckets: 1024 Batches: 1 Memory Usage: 25kB
-> Seq Scan on users u (cost=0.00..12.80 rows=280
width=64) (actual time=0.003..0.017 rows=200 loops=1)
Planning Time: 0.172 ms
Execution Time: 27.916 ms

```

(39 rows)

```
-- Index DDL:
create index idx_records_rating on records (rating);
CREATE INDEX
-- Query against materialized data:
select a.user_id, u.nickname, avg(r.rating) as content_rating
from artists as a
join users as u on a.user_id = u.user_id
join records as r on a.user_id = r.artist_user_id
group by a.user_id, u.nickname
having avg(r.rating) >= 3.3;
```

content_rating	user_id	nickname
3.4566666666666667	user_f921401a-7991-4db2-9491-2cb32b4146db	paynedavid
3.3400000000000000	user_827369a9-7b0e-4937-917e-632d1ed5620f	robertfigueroa

(2 rows)

```
EXPLAIN ANALYZE
select a.user_id, u.nickname, avg(r.rating) as content_rating
from artists as a
join users as u on a.user_id = u.user_id
join records as r on a.user_id = r.artist_user_id
group by a.user_id, u.nickname
having avg(r.rating) >= 3.3;
```

#### QUERY PLAN

```
-----
HashAggregate (cost=100.72..115.72 rows=333 width=96) (actual
time=0.451..0.473 rows=2 loops=1)
  Group Key: a.user_id, u.nickname
  Filter: (avg(r.rating) >= 3.3)
  Batches: 1  Memory Usage: 129kB
  Rows Removed by Filter: 114
  -> Hash Join (cost=40.92..93.22 rows=1000 width=76) (actual
time=0.045..0.297 rows=1000 loops=1)
    Hash Cond: (a.user_id = u.user_id)
    -> Hash Join (cost=24.62..74.27 rows=1000 width=76) (actual
time=0.017..0.165 rows=1000 loops=1)
      Hash Cond: (r.artist_user_id = a.user_id)
      -> Seq Scan on records r (cost=0.00..47.00 rows=1000
width=44) (actual time=0.003..0.031 rows=1000 loops=1)
      -> Hash (cost=16.50..16.50 rows=650 width=32) (actual
time=0.013..0.013 rows=116 loops=1)
        Buckets: 1024  Batches: 1  Memory Usage: 17kB
        -> Seq Scan on artists a (cost=0.00..16.50
```

```

rows=650 width=32) (actual time=0.001..0.006 rows=116 loops=1)
  -> Hash (cost=12.80..12.80 rows=280 width=64) (actual
time=0.027..0.028 rows=200 loops=1)
    Buckets: 1024 Batches: 1 Memory Usage: 25kB
    -> Seq Scan on users u (cost=0.00..12.80 rows=280
width=64) (actual time=0.002..0.014 rows=200 loops=1)
  Planning Time: 0.042 ms
  Execution Time: 0.481 ms
(18 rows)

```

```

/*
... brief performance discussion ...
Without Materialized Data (Using RatedRecords View), the
planning time is 2.533 ms and
execution time is 52.630 ms.

```

```

With Materialized Data and Index (Using Records Table
Directly), the planning time is
1.907 ms and the execution time is only 1.526 ms.

```

```

In conclusion, Using materialized data (with an indexed
rating column in Records) provides
significantly faster and more scalable performance than
dynamically calculating ratings
through the RatedRecords view.

```

```

*/
-- Problem M --
select coalesce(music_quality, 'ALL') as music_quality,
coalesce(device, 'ALL') as device,
count(*) as num_session
from sessions
group by rollup(music_quality, device)
order by num_session desc;

```

music_quality	device	num_session
ALL	ALL	50000
lowest	ALL	8432
lossless	ALL	8400
normal	ALL	8380
low	ALL	8358
high	ALL	8327
Hi-Fi	ALL	8103
lowest	mobile-app	2188
normal	desktop-browser	2156
high	mobile-browser	2153
normal	mobile-app	2147
lowest	desktop-browser	2126
lossless	desktop-app	2115
lossless	desktop-browser	2112
lossless	mobile-app	2106

low	desktop-app	2102
low	mobile-app	2102
Hi-Fi	mobile-app	2100
low	desktop-browser	2091
normal	mobile-browser	2087
lowest	mobile-browser	2075
high	desktop-app	2075
lossless	mobile-browser	2067
low	mobile-browser	2063
high	mobile-app	2051
high	desktop-browser	2048
lowest	desktop-app	2043
Hi-Fi	mobile-browser	2026
Hi-Fi	desktop-app	2003
normal	desktop-app	1990
Hi-Fi	desktop-browser	1974

(31 rows)